

Stack Overflow: Tag Prediction

Assignment 19

abhipise2704@gmail.com_19
(mailto:abhipise2704@gmail.com_19)



```
In [4]: import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import sqlite3
import csv
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from wordcloud import WordCloud
import re
import os
from sqlalchemy import create_engine # database connection
import datetime as dt
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem.snowball import SnowballStemmer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.multiclass import OneVsRestClassifier
from sklearn.linear_model import SGDClassifier
from sklearn import metrics
from sklearn.metrics import f1_score, precision_score, recall_score
from sklearn import svm
from sklearn.linear_model import LogisticRegression
from skmultilearn.adapt import mlknn
from skmultilearn.problem_transform import ClassifierChain
from skmultilearn.problem_transform import BinaryRelevance
from skmultilearn.problem_transform import LabelPowerset
from sklearn.naive_bayes import GaussianNB
from datetime import datetime
from tqdm import tqdm
```

```
In [3]: #!pip install wordcloud
```

```
In [2]: !pip install scikit-multilearn
```

```
Defaulting to user installation because normal site-packages is not writeable
Collecting scikit-multilearn
  Downloading scikit_multilearn-0.2.0-py3-none-any.whl (89 kB)
    |████████████████████████████████████████| 89 kB 3.5 MB/s eta 0:00:011
Installing collected packages: scikit-multilearn
Successfully installed scikit-multilearn-0.2.0
WARNING: You are using pip version 20.0.2; however, version 20.1 is available.
You should consider upgrading via the '/usr/bin/python3 -m pip install --upgrad
e pip' command.
```

```
In [0]: !curl --header 'Host: storage.googleapis.com' --user-agent 'Mozilla/5.0 (Windows
```

% Total	% Received	% Xferd	Average Speed		Time	Time	Time	Current
			Dload	Upload	Total	Spent	Left	Speed
100 2238M	100 2238M	0 0	68.7M	0	0:00:32	0:00:32	--:--:--	93.8M

In [0]: `!unzip /content/Train.zip`

Archive: /content/Train.zip
 inflating: Train.csv

```
In [0]: #https://colab.research.google.com/notebooks/io.ipynb#scrollTo=XDg90BaYqRMd
#https://medium.com/@master_yi/importing-datasets-in-google-colab-c816fc654f97
# Code to read csv file into Colaboratory:
#!pip install -U -q PyDrive
from pydrive.auth import GoogleAuth
from pydrive.drive import GoogleDrive
from google.colab import auth
from oauth2client.client import GoogleCredentials
# Authenticate and create the PyDrive client.
auth.authenticate_user()
gauth = GoogleAuth()
gauth.credentials = GoogleCredentials.get_application_default()
drive = GoogleDrive(gauth)

link = 'https://drive.google.com/open?id=1FQ_Pzk8nzwGq1ZNQ8-FwmTp5wcJy6nx0'

fluff, id = link.split('=')
print (id) # Verify that you have everything after '='

downloaded = drive.CreateFile({'id':id})
downloaded.GetContentFile('train_no_dup.db')
#data = pd.read_csv('preprocessed_data')
# Dataset is now stored in a Pandas Dataframe
#data.shape
```

1FQ_Pzk8nzwGq1ZNQ8-FwmTp5wcJy6nx0

```
In [0]: link = 'https://drive.google.com/open?id=1FQ_Pzk8nzwGq1ZNQ8-FwmTp5wcJy6nx0'

fluff, id = link.split('=')
print (id) # Verify that you have everything after '='

downloaded = drive.CreateFile({'id':id})
downloaded.GetContentFile('Processed.db')
#data = pd.read_csv('preprocessed_data')
# Dataset is now stored in a Pandas Dataframe
#data.shape
```

1FQ_Pzk8nzwGq1ZNQ8-FwmTp5wcJy6nx0

```
In [0]: link = 'https://drive.google.com/open?id=1L5BXggsnc5LmUplfPjkLMfSiP81NAU5'

fluff, id = link.split('=')
print(id) # Verify that you have everything after '='

downloaded = drive.CreateFile({'id':id})
downloaded.GetContentFile('Titlemoreweight.db')
#data = pd.read_csv('preprocessed_data')
# Dataset is now stored in a Pandas Dataframe
#data.shape

1L5BXggsnc5LmUplfPjkLMfSiP81NAU5
```

```
In [0]: !pip install scikit-multilearn
```

```
Collecting scikit-multilearn
  Downloading https://files.pythonhosted.org/packages/bb/1f/e6ff649c72a1cdf2c7a1d31eb21705110ce1c5d3e7e26b2cc300e1637272/scikit\_multilearn-0.2.0-py3-none-any.whl (89kB)
    |████████████████████████████████████████████████████████████████████████████████| 92kB 2.5MB/s
Installing collected packages: scikit-multilearn
Successfully installed scikit-multilearn-0.2.0
```

Stack Overflow: Tag Prediction

1. Business Problem

1.1 Description

Description

Stack Overflow is the largest, most trusted online community for developers to learn, share their programming knowledge, and build their careers.

Stack Overflow is something which every programmer use one way or another. Each month, over 50 million developers come to Stack Overflow to learn, share their knowledge, and build their careers. It features questions and answers on a wide range of topics in computer programming. The website serves as a platform for users to ask and answer questions, and, through membership and active participation, to vote questions and answers up or down and edit questions and answers in a fashion similar to a wiki or Digg. As of April 2014 Stack Overflow has over 4,000,000 registered users, and it exceeded 10,000,000 questions in late August 2015. Based on the type of tags assigned to questions, the top eight most discussed topics on the site are: Java, JavaScript, C#, PHP, Android, jQuery, Python and HTML.

Problem Statemtent

Suggest the tags based on the content that was there in the question posted on Stackoverflow.

Source: <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/>

1.2 Source / useful links

Data Source : <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>
(<https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>)

Youtube : <https://youtu.be/nNDqbUhtIRg> (<https://youtu.be/nNDqbUhtIRg>)

Research paper : <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf> (<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf>)

Research paper : <https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL>
(<https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL>)

1.3 Real World / Business Objectives and Constraints

1. Predict as many tags as possible with high precision and recall.
2. Incorrect tags could impact customer experience on StackOverflow.
3. No strict latency constraints.

2. Machine Learning problem

2.1 Data

2.1.1 Data Overview

Refer: <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>
(<https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>)

All of the data is in 2 files: Train and Test.

Train.csv contains 4 columns: Id,Title,Body,Tags.

Test.csv contains the same columns but without the Tags, which you are to predict.

Size of Train.csv - 6.75GB

Size of Test.csv - 2GB

Number of rows in Train.csv = 6034195

The questions are randomized and contains a mix of verbose text sites as well as sites related to math and programming. The number of questions from each site may vary, and no filtering has been performed on the questions (such as closed questions).

Data Field Explanation

Dataset contains 6,034,195 rows. The columns in the table are:

Id - Unique identifier for each question

Title - The question's title

Body - The body of the question

Tags - The tags associated with the question in a space-separated format (all lowercase, should not contain tabs '\t' or ampersands '&')

2.1.2 Example Data point

Title: Implementing Boundary Value Analysis of Software Testing in a C++ program?

Body :

```

#include<
iostream>\n
#include<
stdlib.h>\n\n
using namespace std;\n\n
int main()\n
{\n
    int n,a[n],x,c,u[n],m[n],e[n][4];\n
    cout<<"Enter the number of variables";\n
cin>>n;\n\n
    cout<<"Enter the Lower, and Upper Limits of the
variables";\n
    for(int y=1; y<n+1; y++)\n
    {\n
        cin>>m[y];\n
        cin>>u[y];\n
    }\n
    for(x=1; x<n+1; x++)\n
    {\n
        a[x] = (m[x] + u[x])/2;\n
    }\n
    c=(n*4)-4;\n
    for(int a1=1; a1<n+1; a1++)\n
    {\n\n
        e[a1][0] = m[a1];\n
        e[a1][1] = m[a1]+1;\n
        e[a1][2] = u[a1]-1;\n
        e[a1][3] = u[a1];\n
    }\n
    for(int i=1; i<n+1; i++)\n
    {\n
        for(int l=1; l<=i; l++)\n
        {\n
            if(l!=1)\n
            {\n
                cout<<a[l]<<"\\t";\n
            }\n
        }\n
        for(int j=0; j<4; j++)\n
        {\n
            cout<<e[i][j];\n
            for(int k=0; k<n-(i+1); k++)\n
            {\n
                cout<<a[k]<<"\\t";\n
            }\n
            cout<<"\\n";\n
        }\n
    }\n
}    \n\n

```

```

        system("PAUSE");\n
        return 0;    \n
    }\n

```

\n\n

<p>The answer should come in the form of a table like</p>\n\n

<pre><code>

```

1          50          50\n
2          50          50\n
99         50          50\n
100        50          50\n
50         1           50\n
50         2           50\n
50         99          50\n
50         100         50\n
50         50          1\n
50         50          2\n
50         50          99\n
50         50          100\n

```

</code></pre>\n\n

<p>if the no of inputs is 3 and their ranges are\n

1,100\n

1,100\n

1,100\n

(could be varied too)</p>\n\n

<p>The output is not coming,can anyone correct the code or tell me what's wrong?</p>\n'

Tags : 'c++ c'

2.2 Mapping the real-world problem to a Machine Learning Problem

2.2.1 Type of Machine Learning Problem

It is a multi-label classification problem

Multi-label Classification: Multilabel classification assigns to each sample a set of target labels. This can be thought as predicting properties of a data-point that are not mutually exclusive, such as topics that are relevant for a document. A question on Stackoverflow might be about any of C, Pointers, FileIO and/or memory-management at the same time or none of these.

__Credit__: <http://scikit-learn.org/stable/modules/multiclass.html>

2.2.2 Performance metric

Micro-Averaged F1-Score (Mean F Score) : The F1 score can be interpreted as a weighted average of the precision and recall, where an F1 score reaches its best value at 1 and worst score at 0. The relative contribution of precision and recall to the F1 score are equal. The formula for the F1 score is:

$$F1 = 2 * (precision * recall) / (precision + recall)$$

In the multi-class and multi-label case, this is the weighted average of the F1 score of each class.

'Micro f1 score':

Calculate metrics globally by counting the total true positives, false negatives and false positives. This is a better metric when we have class imbalance.

'Macro f1 score':

Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

<https://www.kaggle.com/wiki/MeanFScore> (<https://www.kaggle.com/wiki/MeanFScore>)
http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html (http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html)

Hamming loss : The Hamming loss is the fraction of labels that are incorrectly predicted.
<https://www.kaggle.com/wiki/HammingLoss> (<https://www.kaggle.com/wiki/HammingLoss>)

3. Exploratory Data Analysis

3.1 Data Loading and Cleaning

3.1.1 Using Pandas with SQLite to Load the data

In [0]:

```

In [0]: #Creating db file from csv
#Learn SQL: https://www.w3schools.com/sql/default.asp
if not os.path.isfile('train.db'):
    start = datetime.now()
    disk_engine = create_engine('sqlite:///train.db')
    start = dt.datetime.now()
    chunksize = 180000
    j = 0
    index_start = 1
    for df in pd.read_csv('Train.csv', names=['Id', 'Title', 'Body', 'Tags'], chunksize=chunksize):
        df.index += index_start
        j+=1
        print('{} rows'.format(j*chunksize))
        df.to_sql('data', disk_engine, if_exists='append')
        index_start = df.index[-1] + 1
    print("Time taken to run this cell :", datetime.now() - start)

```

3.1.2 Counting the number of rows

```

In [0]: if os.path.isfile('train.db'):
    start = datetime.now()
    con = sqlite3.connect('train.db')
    num_rows = pd.read_sql_query("""SELECT count(*) FROM data""", con)
    #Always remember to close the database
    print("Number of rows in the database :", "\n", num_rows['count(*)'].values[0])
    con.close()
    print("Time taken to count the number of rows :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above cell to

```

Number of rows in the database :

6034196

Time taken to count the number of rows : 0:00:17.687233

3.1.3 Checking for duplicates

```

In [0]: #Learn SQL: https://www.w3schools.com/sql/default.asp
if os.path.isfile('train.db'):
    start = datetime.now()
    con = sqlite3.connect('train.db')
    df_no_dup = pd.read_sql_query('SELECT Title, Body, Tags, COUNT(*) as cnt_dup', con)
    con.close()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the first to generate

```

Time taken to run this cell : 0:05:11.452501

```
In [0]: df_no_dup.head(1)
# we can observe that there are duplicates
```

```
Out[10]:
```

	Title	Body	Tags	cnt_dup
0	Implementing Boundary Value Analysis of S...	<pre>#include<streambuf>\n#include<string>\nusing namespace std;\nint main()\n{\n string s;\n while(s.length() < 100)\n {\n s = s + " ";\n }\n return 0;\n}	c++ c	1

```
In [0]: print("number of duplicate questions :", num_rows['count(*)'].values[0] - df_no_dup['cnt_dup'].sum())

number of duplicate questions : 1827881 ( 30.292038906260256 % )
```

```
In [0]: # number of times each question appeared in our database
df_no_dup.cnt_dup.value_counts()
```

```
Out[12]:
```

1	2656284
2	1272336
3	277575
4	90
5	25
6	5

Name: cnt_dup, dtype: int64

```
In [0]: #https://stackoverflow.com/questions/39092067/pandas-dataframe-convert-column-type
df_no_dup["Tags"] = df_no_dup.Tags.astype(str)
```

```
In [0]: start = datetime.now()
df_no_dup["tag_count"] = df_no_dup["Tags"].apply(lambda text: len(text.split(" ")))
# adding a new feature number of tags per question
print("Time taken to run this cell :", datetime.now() - start)
df_no_dup.head()
```

Time taken to run this cell : 0:00:02.798297

```
Out[16]:
```

	Title	Body	Tags	cnt_dup	tag_
0	Implementing Boundary Value Analysis of S...	<pre>#include<streambuf>\n#include<string>\nusing namespace std;\nint main()\n{\n string s;\n while(s.length() < 100)\n {\n s = s + " ";\n }\n return 0;\n}	c++ c	1	
1	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding	1	
2	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding columns	1	
3	java.lang.NoClassDefFoundError: javax/serv...	<p>I followed the guide in <a href="http://sta...	jsp jstl	1	
4	java.sql.SQLException: [Microsoft] [ODBC Dri...	<p>I use the following code</p>\n\n<pre>\n\n</pre>	java jdbc	2	

```
In [0]: # distribution of number of tags per question
df_no_dup.tag_count.value_counts()
```

```
Out[17]: 3    1206157
         2    1111706
         4     814996
         1     568298
         5     505158
         Name: tag_count, dtype: int64
```

```
In [0]: con = sqlite3.connect('/content/train_no_dup.sqlite')
```

```
In [ ]: #Creating a new database with no duplicates
if not os.path.isfile('/content/train_no_dup.db'):
    disk_dup = create_engine("sqlite:///content/train_no_dup.db")
    no_dup = pd.DataFrame(df_no_dup, columns=['Title', 'Body', 'Tags'])
    no_dup.to_sql('no_dup_train', disk_dup)
```

```
In [ ]: #This method seems more appropriate to work with this much data.
#creating the connection with database file.
if os.path.isfile('/content/train_no_dup.db'):
    start = datetime.now()
    con = sqlite3.connect('/content/train_no_dup.sqlite')
    tag_data = pd.read_sql_query("""SELECT Tags FROM train_no_dup """, con)
    #Always remember to close the database
    con.close()

    # Let's now drop unwanted column.
    tag_data.drop(tag_data.index[0], inplace=True)
    #Printing first 5 columns from our data frame
    tag_data.head()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above cells to")
```

```
In [0]:
```

3.2 Analysis of Tags

3.2.1 Total number of unique tags

```
In [0]: # Importing & Initializing the "CountVectorizer" object, which
        #is scikit-learn's bag of words tool.

        #by default 'split()' will tokenize each tag using space.
        vectorizer = CountVectorizer(tokenizer = lambda x: x.split())
        # fit_transform() does two functions: First, it fits the model
        # and learns the vocabulary; second, it transforms our training data
        # into feature vectors. The input to fit_transform should be a list of strings.
        tag_dtm = vectorizer.fit_transform(tag_data['Tags'])
```

```
In [0]: print("Number of data points :", tag_dtm.shape[0])
        print("Number of unique tags :", tag_dtm.shape[1])
```

Number of data points : 4206314
Number of unique tags : 42048

```
In [0]: # 'get_feature_name()' gives us the vocabulary.
        tags = vectorizer.get_feature_names()
        # Lets look at the tags we have.
        print("Some of the tags we have :", tags[:10])
```

Some of the tags we have : ['.a', '.app', '.asp.net-mvc', '.aspxauth', '.bash-profile', '.class-file', '.cs-file', '.doc', '.drv', '.ds-store']

```
In [0]: #import dill
        #dill.dump_session('stack_overflow_1.db')
```

3.2.3 Number of times a tag appeared

```
In [0]: # https://stackoverflow.com/questions/15115765/how-to-access-sparse-matrix-elements
        # Lets now store the document term matrix in a dictionary.
        freqs = tag_dtm.sum(axis=0).A1
        result = dict(zip(tags, freqs))
```


```
In [0]: #Saving this dictionary to csv files.
if not os.path.isfile('tag_counts_dict_dtm.csv'):
    with open('tag_counts_dict_dtm.csv', 'w') as csv_file:
        writer = csv.writer(csv_file)
        for key, value in tqdm(result.items()):
            writer.writerow([key, value])
tag_df = pd.read_csv("tag_counts_dict_dtm.csv", names=['Tags', 'Counts'])
tag_df.head()
```

```
Out[17]:
```


	Tags	Counts
0	.a	18
1	.app	37
2	.asp.net-mvc	1
3	.aspxauth	21
4	.bash-profile	138

```
In [0]: tag_df_sorted = tag_df.sort_values(['Counts'], ascending=False)
tag_counts = tag_df_sorted['Counts'].values
```

```
In [0]: plt.plot(tag_counts)
plt.title("Distribution of number of times tag appeared questions")
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
```




```
In [0]: plt.plot(tag_counts[0:10000])
plt.title('first 10k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:10000:25]), tag_counts[0:10000:25])
```



400	[331505	44829	22429	17728	13364	11162	10029	9148	8054	7151
6466	5865	5370	4983	4526	4281	4144	3929	3750	3593	
3453	3299	3123	2989	2891	2738	2647	2527	2431	2331	
2259	2186	2097	2020	1959	1900	1828	1770	1723	1673	
1631	1574	1532	1479	1448	1406	1365	1328	1300	1266	
1245	1222	1197	1181	1158	1139	1121	1101	1076	1056	
1038	1023	1006	983	966	952	938	926	911	891	
882	869	856	841	830	816	804	789	779	770	
752	743	733	725	712	702	688	678	671	658	
650	643	634	627	616	607	598	589	583	577	
568	559	552	545	540	533	526	518	512	506	
500	495	490	485	480	477	469	465	457	450	
447	442	437	432	426	422	418	413	408	403	
398	393	388	385	381	378	374	370	367	365	
361	357	354	350	347	344	342	339	336	332	
330	326	323	319	315	312	309	307	304	301	
299	296	293	291	289	286	284	281	278	276	
275	272	270	268	265	262	260	258	256	254	
252	250	249	247	245	243	241	239	238	236	
234	233	232	230	228	226	224	222	220	219	
217	215	214	212	210	209	207	205	204	203	
201	200	199	198	196	194	193	192	191	189	


188	186	185	183	182	181	180	179	178	177
175	174	172	171	170	169	168	167	166	165
164	162	161	160	159	158	157	156	156	155
154	153	152	151	150	149	149	148	147	146
145	144	143	142	142	141	140	139	138	137
137	136	135	134	134	133	132	131	130	130
129	128	128	127	126	126	125	124	124	123
123	122	122	121	120	120	119	118	118	117
117	116	116	115	115	114	113	113	112	111
111	110	109	109	108	108	107	106	106	106
105	105	104	104	103	103	102	102	101	101
100	100	99	99	98	98	97	97	96	96
95	95	94	94	93	93	93	92	92	91
91	90	90	89	89	88	88	87	87	86
86	86	85	85	84	84	83	83	83	82
82	82	81	81	80	80	80	79	79	78
78	78	78	77	77	76	76	76	75	75
75	74	74	74	73	73	73	73	72	72]


```
In [0]: plt.plot(tag_counts[0:1000])
plt.title('first 1k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:1000:5]), tag_counts[0:1000:5])
```



```
200 [331505 221533 122769 95160 62023 44829 37170 31897 26925 24537
22429 21820 20957 19758 18905 17728 15533 15097 14884 13703
13364 13157 12407 11658 11228 11162 10863 10600 10350 10224
10029 9884 9719 9411 9252 9148 9040 8617 8361 8163
8054 7867 7702 7564 7274 7151 7052 6847 6656 6553
6466 6291 6183 6093 5971 5865 5760 5577 5490 5411
5370 5283 5207 5107 5066 4983 4891 4785 4658 4549
4526 4487 4429 4335 4310 4281 4239 4228 4195 4159
4144 4088 4050 4002 3957 3929 3874 3849 3818 3797
3750 3703 3685 3658 3615 3593 3564 3521 3505 3483
3453 3427 3396 3363 3326 3299 3272 3232 3196 3168
3123 3094 3073 3050 3012 2989 2984 2953 2934 2903
2891 2844 2819 2784 2754 2738 2726 2708 2681 2669
2647 2621 2604 2594 2556 2527 2510 2482 2460 2444
2431 2409 2395 2380 2363 2331 2312 2297 2290 2281
2259 2246 2222 2211 2198 2186 2162 2142 2132 2107
2097 2078 2057 2045 2036 2020 2011 1994 1971 1965
1959 1952 1940 1932 1912 1900 1879 1865 1855 1841
1828 1821 1813 1801 1782 1770 1760 1747 1741 1734
1723 1707 1697 1688 1683 1673 1665 1656 1646 1639]
```

```
In [0]: plt.plot(tag_counts[0:500])
plt.title('first 500 tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:500:5]), tag_counts[0:500:5])
```



```
100 [331505 221533 122769 95160 62023 44829 37170 31897 26925 24537
22429 21820 20957 19758 18905 17728 15533 15097 14884 13703
13364 13157 12407 11658 11228 11162 10863 10600 10350 10224
10029 9884 9719 9411 9252 9148 9040 8617 8361 8163
8054 7867 7702 7564 7274 7151 7052 6847 6656 6553
6466 6291 6183 6093 5971 5865 5760 5577 5490 5411
5370 5283 5207 5107 5066 4983 4891 4785 4658 4549
4526 4487 4429 4335 4310 4281 4239 4228 4195 4159
4144 4088 4050 4002 3957 3929 3874 3849 3818 3797
3750 3703 3685 3658 3615 3593 3564 3521 3505 3483]
```


```

In [0]: plt.plot(tag_counts[0:100], c='b')
plt.scatter(x=list(range(0,100,5)), y=tag_counts[0:100:5], c='orange', label="quar
# quantiles with 0.25 difference
plt.scatter(x=list(range(0,100,25)), y=tag_counts[0:100:25], c='m', label = "quar

for x,y in zip(list(range(0,100,25)), tag_counts[0:100:25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500))

plt.title('first 100 tags: Distribution of number of times tag appeared questions
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.legend()
plt.show()
print(len(tag_counts[0:100:5]), tag_counts[0:100:5])

```



```

20 [331505 221533 122769 95160 62023 44829 37170 31897 26925 24537
    22429 21820 20957 19758 18905 17728 15533 15097 14884 13703]

```

```
In [0]: # Store tags greater than 10K in one List
lst_tags_gt_10k = tag_df[tag_df.Counts>10000].Tags
#Print the Length of the List
print ('{} Tags are used more than 10000 times'.format(len(lst_tags_gt_10k)))
# Store tags greater than 100K in one List
lst_tags_gt_100k = tag_df[tag_df.Counts>100000].Tags
#Print the Length of the List.
print ('{} Tags are used more than 100000 times'.format(len(lst_tags_gt_100k)))
```

153 Tags are used more than 10000 times
 14 Tags are used more than 100000 times

Observations:

1. There are total 153 tags which are used more than 10000 times.
2. 14 tags are used more than 100000 times.
3. Most frequent tag (i.e. c#) is used 331505 times.
4. Since some tags occur much more frequently than others, Micro-averaged F1-score is the appropriate metric for this problem.

3.2.4 Tags Per Question

```
In [0]: #Storing the count of tag in each question in list 'tag_count'
tag_quest_count = tag_dtm.sum(axis=1).tolist()
#Converting List of Lists into single List, we will get [[3], [4], [2], [2], [3]]
tag_quest_count=[int(j) for i in tag_quest_count for j in i]
print ('We have total {} datapoints.'.format(len(tag_quest_count)))


print(tag_quest_count[:5])
```

We have total 4206314 datapoints.
 [3, 4, 2, 2, 3]

```
In [0]: print( "Maximum number of tags per question: %d"%max(tag_quest_count))
print( "Minimum number of tags per question: %d"%min(tag_quest_count))
print( "Avg. number of tags per question: %f"% ((sum(tag_quest_count)*1.0)/len(tag_quest_count)))
```

Maximum number of tags per question: 5
 Minimum number of tags per question: 1
 Avg. number of tags per question: 2.899440

```
In [0]: sns.countplot(tag_quest_count, palette='gist_rainbow')
plt.title("Number of tags in the questions ")
plt.xlabel("Number of Tags")
plt.ylabel("Number of questions")
plt.show()
```



Observations:


1. Maximum number of tags per question: 5
2. Minimum number of tags per question: 1
3. Avg. number of tags per question: 2.899
4. Most of the questions are having 2 or 3 tags

3.2.5 Most Frequent Tags

```
In [0]: # Plotting word cloud
start = datetime.now()

# Lets first convert the 'result' dictionary to 'list of tuples'
tup = dict(result.items())
#Initializing WordCloud using frequencies of tags.
wordcloud = WordCloud(    background_color='black',
                           width=1600,
                           height=800,
                           ).generate_from_frequencies(tup)

fig = plt.figure(figsize=(30,20))
plt.imshow(wordcloud)
plt.axis('off')
plt.tight_layout(pad=0)
fig.savefig("tag.png")
plt.show()
print("Time taken to run this cell :", datetime.now() - start)
```




Time taken to run this cell : 0:00:09.527878

Observations:

A look at the word cloud shows that "c#", "java", "php", "asp.net", "javascript", "c++" are some of the most frequent tags.

3.2.6 The top 20 tags

```
In [0]: i=np.arange(30)
tag_df_sorted.head(30).plot(kind='bar')
plt.title('Frequency of top 20 tags')
plt.xticks(i, tag_df_sorted['Tags'])
plt.xlabel('Tags')
plt.ylabel('Counts')
plt.show()
```



Observations:

1. Majority of the most frequent tags are programming language.
2. C# is the top most frequent programming language.
3. Android, IOS, Linux and windows are among the top most frequent operating systems.

3.3 Cleaning and preprocessing of Questions

3.3.1 Preprocessing

1. Sample 1M data points
2. Separate out code-snippets from Body
3. Remove Special characters from Question title and description (not in code)
4. Remove stop words (Except 'C')
5. Remove HTML Tags
6. Convert all the characters into small letters
7. Use SnowballStemmer to stem the words

```
In [7]: import nltk
nltk.download('stopwords')

def striphtml(data):
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', str(data))
    return cleantext
stop_words = set(stopwords.words('english'))
stemmer = SnowballStemmer("english")

[nltk_data] Downloading package stopwords to
[nltk_data]   /home/abhipise2704/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
```



```
In [6]: #http://www.sqlitetutorial.net/sqlite-python/create-tables/
def create_connection(db_file):
    """ create a database connection to the SQLite database
        specified by db_file
    :param db_file: database file
    :return: Connection object or None
    """
    try:
        conn = sqlite3.connect(db_file)
        return conn
    except Error as e:
        print(e)

    return None

def create_table(conn, create_table_sql):
    """ create a table from the create_table_sql statement
    :param conn: Connection object
    :param create_table_sql: a CREATE TABLE statement
    :return:
    """
    try:
        c = conn.cursor()
        c.execute(create_table_sql)
    except Error as e:
        print(e)

def checkTableExists(dbcon):
    cursr = dbcon.cursor()
    str = "select name from sqlite_master where type='table'"
    table_names = cursr.execute(str)
    print("Tables in the databse:")
    tables = table_names.fetchall()
    print(tables[0][0])
    return(len(tables))

def create_database_table(database, query):
    conn = create_connection(database)
    if conn is not None:
        create_table(conn, query)
        checkTableExists(conn)
    else:
        print("Error! cannot create the database connection.")
    conn.close()
```

```
In [0]: #1000000
```

```
In [0]: # http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table
start = datetime.now()
read_db = 'train_no_dup.db'
write_db = 'Processed.db'
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RAND()")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")
print("Time taken to run this cell :", datetime.now() - start)
```

__ we create a new data base to store the sampled and preprocessed questions __

In [0]: [#http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/](http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/)

```
start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_proccesed = 0
for row in tqdm(reader):

    is_code = 0

    title, question, tags = row[0], row[1], row[2]

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=striphtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    question=str(title)+" "+str(question)
    question=re.sub(r'^[A-Za-z]+', ' ',question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt for the l
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_words

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,v
    if (questions_proccesed%100000==0):
        print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_le
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_le
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/q

print("Time taken to run this cell :", datetime.now() - start)
```

```
number of questions completed= 100000
number of questions completed= 200000
number of questions completed= 300000
number of questions completed= 400000
number of questions completed= 500000
```

```
number of questions completed= 600000  
number of questions completed= 700000  
number of questions completed= 800000  
number of questions completed= 900000  
Avg. length of questions(Title+Body) before processing: 1169  
Avg. length of questions(Title+Body) after processing: 327  
Percent of questions containing code: 57  
Time taken to run this cell : 0:47:05.946582
```

```
In [0]: # dont forget to close the connections, or else you will end up with Locks  
conn_r.commit()  
conn_w.commit()  
conn_r.close()  
conn_w.close()
```

```
In [0]: if os.path.isfile(write_db):
        conn_r = create_connection(write_db)
        if conn_r is not None:
            reader = conn_r.cursor()
            reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
            print("Questions after preprocessed")
            print('='*100)
            reader.fetchone()
            for row in reader:
                print(row)
                print('-'*100)
        conn_r.commit()
        conn_r.close()
```

Questions after preprocessed

```
=====
=====
('ef code first defin one mani relationship differ key troubl defin one zero
mani relationship entiti ef object model look like use fluent api object comp
osit pk defin batch id batch detail id use fluent api object composi pk defi
n batch detail id compani id map exist databas tpt basic idea submittedtransa
ct zero mani submittedsplitttransact associ navig realli need one way submitte
dtransact submittedsplitttransact need dbcontext class onmodelcr overrid map c
lass lazi load occur submittedtransact submittedsplitttransact help would much
appreci edit taken advic made follow chang dbcontext class ad follow onmodelc
r overrid must miss someth get follow except thrown submittedtransact key bat
ch id batch detail id zero one mani submittedsplitttransact key batch detail i
d compani id rather assum convent creat relationship two object configur requ
ir sinc obvious wrong',)
-----
-----
('explan new statement review section c code came accross statement block com
e accross new oper use way someon explain new call way',)
-----
-----
('error function notat function solv logic riddl iloczyni list structur list
possibl candid solut list possibl coordin matrix wan na choos one candid comp
ar possibl candid element equal wan na delet coordin call function skasuj loo
k like ni knowledg haskel cant see what wrong',)
-----
-----
('step plan move one isp anoth one work busi plan switch isp realli soon need
chang lot inform dns wan wan wifi question guy help mayb peopl plan correct c
hang current isp new one first dns know receiv new ip isp major chang need ta
ke consider exchang server owa vpn two site link wireless connect km away cit
rix server vmware exchang domain control link place import server crucial ste
p inform need know avoid downtim busi regard ndavid',)
-----
-----
('use ef migrat creat databas googl migrat tutori af first run applic creat d
atabas ef enabl migrat way creat databas migrat rune applic tri',)
-----
-----
('magento unit test problem magento site recent look way check integr magento
site given point unit test jump one method would assum would big job write wh
ole lot test check everyth site work anyon involv unit test magento advis fol
low possibl test whole site custom modul nis exampl test would amaz given sit
```

e heavili link databas would nbe possibl fulli test site without disturb data
bas better way automaticlli check integr magento site say integr realli mean
fault site ship payment etc work correct',)

('find network devic without bonjour write mac applic need discov mac pcs iph
on ipad connect wifi network bonjour seem reason choic turn problem mani type
router mine exampl work block bonjour servic need find ip devic tri connect a
pplic specif port determin process run best approach accomplish task without
violat app store sandbox',)

('send multipl row mysql databas want send user mysql databas column user ski
ll time nnow want abl add one row user differ time etc would code send databa
s nthen use help schema',)

('insert data mysql php powerpoint event powerpoint present run continu way u
pdat slide present automat data mysql databas websit',)

```
In [0]: #Taking 1 Million entries to a dataframe.
write_db = 'Processed.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM Quest
conn_r.commit()
conn_r.close()
```

```
In [0]: preprocessed_data.head()
```

```
Out[47]:
```

	question	tags
0	resiz root window tkinter	python tkinter
1	ef code first defin one mani relationship diff...	entity-framework-4.1
2	explan new statement review section c code cam...	c++
3	error function notat function solv logic riddl...	haskell logic
4	step plan move one isp anoth one work busi pla...	dns isp

```
In [0]: print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

```
number of data points in sample : 999999
number of dimensions : 2
```

4. Machine Learning Models

4.1 Converting tags for multilabel problems

X	y1	y2	y3	y4
x1	0	1	1	0
x1	1	0	0	0
x1	0	1	0	0

```
In [0]: # binary='true' will give a binary vectorizer
vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```


__ We will sample the number of tags instead considering all of them (due to limitation of computing power) __

```
In [19]: def tags_to_choose(n):
    t = multilabel_y.sum(axis=0).tolist()[0]
    sorted_tags_i = sorted(range(len(t)), key=lambda i: t[i], reverse=True)
    multilabel_yn=multilabel_y[:,sorted_tags_i[:n]]
    return multilabel_yn

def questions_explained_fn(n):
    multilabel_yn = tags_to_choose(n)
    x= multilabel_yn.sum(axis=1)
    return (np.count_nonzero(x!=0))
```

```
In [0]: questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/tot
```

```
In [0]: fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()
# you can choose any number of tags based on your computing power, minimum is 500
print("with ", 5500, "tags we are covering ", questions_explained[50], "% of questions")
```



with 5500 tags we are covering 99.04 % of questions

```
In [0]: multilabel_yx = tags_to_choose(5500)
print("number of questions that are not covered :", questions_explained_fn(5500),
      "\nnumber of questions that are not covered : 9599 out of 999999")
```

```
In [0]: print("Number of tags in sample :", multilabel_y.shape[1])
print("number of tags taken :", multilabel_yx.shape[1], "(", (multilabel_yx.shape[1] / multilabel_y.shape[1]) * 100, "%)")

Number of tags in sample : 35422
number of tags taken : 5500 ( 15.527073570097679 %)
```

__ We consider top 15% tags which covers 99% of the questions __

4.2 Split the data into test and train (80:20)

```
In [0]: total_size=preprocessed_data.shape[0]
train_size=int(0.80*total_size)

x_train=preprocessed_data.head(train_size)
x_test=preprocessed_data.tail(total_size - train_size)

y_train = multilabel_yx[0:train_size,:]
y_test = multilabel_yx[train_size:total_size,:]
```



```
In [0]: print("Number of data points in train data :", y_train.shape)
        print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (799999, 5500)

Number of data points in test data : (200000, 5500)

4.3 Featurizing data

Type *Markdown* and LaTeX: α^2

```
In [0]: start = datetime.now()
        vectorizer = CountVectorizer(min_df=0.00009, max_features=200000,
                                     tokenizer = lambda x: x.split(),
                                     ngram_range=(1,4))

        x_train_multilabel = vectorizer.fit_transform(x_train['question'])
        x_test_multilabel = vectorizer.transform(x_test['question'])
        print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:09:50.460431

```
In [0]: print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shape)
        print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)
```

Diamensions of train data X: (799999, 88244) Y : (799999, 5500)

Diamensions of test data X: (200000, 88244) Y: (200000, 5500)

```
In [0]: # https://www.analyticsvidhya.com/blog/2017/08/introduction-to-multi-label-classification/
#https://stats.stackexchange.com/questions/117796/scikit-multi-label-classification
# classifier = LabelPowerset(GaussianNB())
"""
from skmultilearn.adapt import MLkNN
classifier = MLkNN(k=21)

# train
classifier.fit(x_train_multilabel, y_train)

# predict
predictions = classifier.predict(x_test_multilabel)
print(accuracy_score(y_test, predictions))
print(metrics.f1_score(y_test, predictions, average = 'macro'))
print(metrics.f1_score(y_test, predictions, average = 'micro'))
print(metrics.hamming_loss(y_test, predictions))

"""
# we are getting memory error because the multilearn package
# is trying to convert the data into dense matrix
# -----
#MemoryError                                Traceback (most recent call last)
#<ipython-input-170-f0e7c7f3e0be> in <module>()
#----> classifier.fit(x_train_multilabel, y_train)
```

```
Out[92]: "\nfrom skmultilearn.adapt import MLkNN\nnclassifier = MLkNN(k=21)\n\n# train\nclassifier.fit(x_train_multilabel, y_train)\n\n# predict\npredictions = classifier.predict(x_test_multilabel)\nprint(accuracy_score(y_test, predictions))\nprint(metrics.f1_score(y_test, predictions, average = 'macro'))\nprint(metrics.f1_score(y_test, predictions, average = 'micro'))\nprint(metrics.hamming_loss(y_test, predictions))\n\n"
```

4.4 Applying Logistic Regression with OneVsRest Classifier

```
In [0]: # this will be taking so much time try not to run it, download the lr_with_equal_
# This takes about 6-7 hours to run.
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)

print("accuracy :",metrics.accuracy_score(y_test,predictions))
print("macro f1 score :",metrics.f1_score(y_test, predictions, average = 'macro')
print("micro f1 scoore :",metrics.f1_score(y_test, predictions, average = 'micro')
print("hamming loss :",metrics.hamming_loss(y_test,predictions))
print("Precision recall report :\n",metrics.classification_report(y_test, predict
```

```
accuracy : 0.081965
macro f1 score : 0.0963020140154
micro f1 scoore : 0.374270748817
hamming loss : 0.00041225090909090907
Precision recall report :
              precision    recall  f1-score   support

     0         0.62         0.23         0.33         15760
     1         0.79         0.43         0.56         14039
     2         0.82         0.55         0.66         13446
     3         0.76         0.42         0.54         12730
     4         0.94         0.76         0.84         11229
     5         0.85         0.64         0.73         10561
     6         0.70         0.30         0.42          6958
     7         0.87         0.61         0.72          6309
     8         0.70         0.40         0.50          6032
     9         0.78         0.43         0.55          6020
    10         0.86         0.62         0.72          5707
    11         0.52         0.17         0.25          5723
    12         0.55         0.10         0.16          5531
```

```
In [0]: from sklearn.externals import joblib
joblib.dump(classifier, 'lr_with_equal_weight.pkl')
```

4.5 Modeling with less data points (0.2M data points) and more weight to title and 500 tags only.

```
In [5]: !curl --header 'Host: doc-0s-9c-docs.googleusercontent.com' --user-agent 'Mozilla
```

```
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed
100 5878M    0 5878M    0     0  68.6M      0  --:--:--  0:01:25 --:--:-- 150M
```

```
In [8]: sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text, title text, body text, tags text)
create_database_table("Titlmoreweight.db", sql_create_table)
```

Tables in the database:
QuestionsProcessed

```
In [9]: # http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table

read_db = 'train_no_dup.db'
write_db = 'Titlmoreweight.db'
train_datasize = 160000
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        # for selecting first 0.5M rows
        reader.execute("SELECT Title, Body, Tags From no_dup_train LIMIT 200001;")
        # for selecting random points
        #reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RAND()")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")
```

Tables in the database:
QuestionsProcessed
Cleared All the rows

4.5.1 Preprocessing of questions

1. Separate Code from Body
2. Remove Special characters from Question title and description (not in code)
3. **Give more weightage to title : Add title three times to the question**

```
<li> Remove stop words (Except 'C') </li>
<li> Remove HTML Tags </li>
<li> Convert all the characters into small letters </li>
<li> Use SnowballStemmer to stem the words </li>
```

```

In [10]: #http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/
start = datetime.now()

import nltk
nltk.download('punkt')
from tqdm import tqdm

preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_proccesed = 0
for row in tqdm(reader):

    is_code = 0

    title, question, tags = row[0], row[1], str(row[2])

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=stripthtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    # adding title three time to the data to increase its weight
    # add tags string to the training data

    question=str(title)+" "+str(title)+" "+str(title)+" "+question

    # if questions_proccesed<=train_datasize:
    #     question=str(title)+" "+str(title)+" "+str(title)+" "+question+" "+str(tags)
    # else:
    #     question=str(title)+" "+str(title)+" "+str(title)+" "+question

    question=re.sub(r'^[A-Za-z0-9#+.\-]+', ' ', question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt for the letter 'l'
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_words)

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,words_post,length,is_code) values(?,?,?,?,?,?,?)")
    if (questions_proccesed%100000==0):
        print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed

```

```

no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_len_post)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/questions_proccesed))

print("Time taken to run this cell :", datetime.now() - start)

```

```

[nltk_data] Downloading package punkt to
[nltk_data]      /home/abhipise2704/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
100037it [03:43, 417.61it/s]

```

```

number of questions completed= 100000

```

```

200000it [07:15, 458.90it/s]

```

```

number of questions completed= 200000

```

```

Avg. length of questions(Title+Body) before processing: 1322

```

```

Avg. length of questions(Title+Body) after processing: 429

```

```

Percent of questions containing code: 57

```

```

Time taken to run this cell : 0:07:17.429439

```

```

In [ ]: #

```

```

In [11]: # never forget to close the conections or else we will end up with database locks
conn_r.commit()
conn_w.commit()
conn_r.close()
conn_w.close()

```

```

In [12]: #import dill
#dill.dump_session('stack_overflow_11.db')

```

___ Sample quesitons after preprocessing of data ___

```
In [13]: if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
        print("Questions after preprocessed")
        print('='*100)
        reader.fetchone()
        for row in reader:
            print(row)
            print('-'*100)
conn_r.commit()
conn_r.close()
```

Questions after preprocessed

```
=====
=====
('dynam datagrid bind silverlight dynam datagrid bind silverlight dynam datagri
d bind silverlight bind datagrid dynam code wrote code debug code block seem bi
nd correct grid come column form come grid column although necessari bind nthan
k repli advance..',)
-----
-----
('java.lang.noclassdeffounderror javax servlet jsp tagext taglibraryvalid java.
lang.noclassdeffounderror javax servlet jsp tagext taglibraryvalid java.lang.no
classdeffounderror javax servlet jsp tagext taglibraryvalid follow guid link in
stal jstl got follow error tri launch jsp page java.lang.noclassdeffounderror j
avax servlet jsp tagext taglibraryvalid taglib declar instal jstl 1.1 tomcat we
bapp tri project work also tri version 1.2 jstl still messag caus solv',)
-----
-----
('java.sql.sqlexcept microsoft odbc driver manag invalid descriptor index java.
sql.sqlexcept microsoft odbc driver manag invalid descriptor index java.sql.sql
except microsoft odbc driver manag invalid descriptor index use follow code dis
play caus solv',)
-----
-----
('better way updat feed fb php sdk better way updat feed fb php sdk better way
updat feed fb php sdk novic facebook api read mani tutori still confused.i find
post feed api method like correct second way use curl someth like way better',)
-----
-----
('btnadd click event open two window record ad btnadd click event open two wind
ow record ad btnadd click event open two window record ad open window search.as
px use code hav add button search.aspx nwhen insert record btnadd click event o
pen anoth window nafter insert record close window',)
-----
-----
('sql inject issu prevent correct form submiss php sql inject issu prevent corr
ect form submiss php sql inject issu prevent correct form submiss php check eve
ryth think make sure input field safe type sql inject good news safe bad news o
ne tag mess form submiss place even touch life figur exact html use templat fil
e forgiv okay entir php script get execut see data post none forum field post p
roblem use someth titl field none data get post current use print post see subm
it noth work flawless statement though also mention script work flawless local
machin use host come across problem state list input test mess',)
-----
```

```
-----
('countabl subaddit lebesgu measur countabl subaddit lebesgu measur countabl su
baddit lebesgu measur let lbrace rbrace sequenc set sigma -algebra mathcal want
show left bigcup right leq sum left right countabl addit measur defin set sigma
algebra mathcal think use monoton properti somewher proof start appreci littl h
elp nthank ad han answer make follow addit construct given han answer clear big
cup bigcup cap emptyset neq left bigcup right left bigcup right sum left right
also construct subset monoton left right leq left right final would sum leq sum
result follow',)
-----
```

```
-----
('hql equival sql queri hql equival sql queri hql equival sql queri hql queri r
eplac name class properti name error occur hql error',)
-----
```

```
-----
('undefin symbol architectur i386 objc class skpsmtpmessag referenc error undef
in symbol architectur i386 objc class skpsmtpmessag referenc error undefin symb
ol architectur i386 objc class skpsmtpmessag referenc error import framework se
nd email applic background import framework i.e skpsmtpmessag somebodi suggest
get error collect2 ld return exit status import framework correct sorc taken fr
amework follow mfmcomposeviewcontrol question lock field updat answer drag d
rop folder project click copi nthat',)
-----
```

__ Saving Preprocessed data to a Database __

```
In [14]: #Taking 0.5 Million entries to a dataframe.
write_db = 'Titlemoreweight.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM Quest
conn_r.commit()
conn_r.close()
```

```
In [15]: preprocessed_data.head()
```

```
Out[15]:
```

	question	tags
0	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding
1	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding columns
2	java.lang.noclassdeffoundererror javax servlet j...	jsp jstl
3	java.sql.sqlexcept microsoft odbc driver manag...	java jdbc
4	better way updat feed fb php sdk better way up...	facebook api facebook-php-sdk

```
In [16]: print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

```
number of data points in sample : 200000
number of dimensions : 2
```


__ Converting string Tags to multilable output variables __


```
In [17]: vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```

__ Selecting 500 Tags __

```
In [20]: questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/tot
```

```
In [21]: fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()
# you can choose any number of tags based on your computing power, minimun is 500
print("with ",5500,"tags we are covering ",questions_explained[50],"% of question
print("with ",500,"tags we are covering ",questions_explained[0],"% of questions'
```



with 5500 tags we are covering 99.41 % of questions
 with 500 tags we are covering 92.478 % of questions

```
In [22]: # we will be taking 500 tags
multilabel_yx = tags_to_choose(500)
print("number of questions that are not covered :", questions_explained_fn(500),'
```

number of questions that are not covered : 15044 out of 200000

```
In [23]: x_train=preprocessed_data.head(train_datasize)
x_test=preprocessed_data.tail(preprocessed_data.shape[0] - 160000)

y_train = multilabel_yx[0:train_datasize,:]
y_test = multilabel_yx[train_datasize:preprocessed_data.shape[0],:]
```

```
In [24]: print("Number of data points in train data :", y_train.shape)
print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (160000, 500)

Number of data points in test data : (40000, 500)

ASSIGNMENT STACK OVERFLOW TAGGING

1. Use bag of words upto 4 grams and compute the micro f1 score with Logistic regression(OvR)
2. Perform hyperparam tuning on alpha (or lambda) for Logistic regression to improve the performance using GridSearch
3. Try OneVsRestClassifier with Linear-SVM (SGDClassifier with loss-hinge)

1: Assignment Task 1---->> Use bag of words upto 4 grams and compute the micro f1 score with Logistic regression(OvR)

1.2: Featurizing data with Count vectorizer(BOW) For 0.2 Million Datapoints and 500 labels

```
In [25]: start = datetime.now()
vectorizer = CountVectorizer(min_df=0.00009, max_features=10000,
                             tokenizer = lambda x: x.split(),
                             ngram_range=(1,4))

x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:03:23.797788

```
In [26]: print("Dimensions of train data X Bow :",x_train_multilabel.shape, "Y :",y_train.shape)
print("Dimensions of test data X bow :",x_test_multilabel.shape,"Y:",y_test.shape)
```

Dimensions of train data X Bow : (160000, 10000) Y : (160000, 500)

Dimensions of test data X bow : (40000, 10000) Y: (40000, 500)

2: Assignment Task 2---->> Perform hyperparam tuning on alpha (or lambda) for Logistic

regression to improve the performance using GridSearch

2.1: GRIDSEARCHCV TO FIND HYPERPARAMETER TUNING ALPHA

```

In [28]: #https://stackoverflow.com/questions/56084591/how-to-do-gridsearchcv-for-f1-score
##https://stackoverflow.com/questions/56416576/getting-keyerror-from-sklearn-model
# https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridS

from datetime import datetime
start = datetime.now()
from sklearn.model_selection import GridSearchCV
from scipy.stats import randint as sp_randint
from sklearn.model_selection import RandomizedSearchCV
#from sklearn.model_selection import GridSearchCV
import math
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import f1_score, make_scorer

#f1 = make_scorer(f1_score , average='micro')

classifier_LR_tune = OneVsRestClassifier(SGDClassifier(loss='log',penalty='l1'))
parameters = {'estimator__alpha':[0.00001,0.0001,0.001,0.01]}

clf = GridSearchCV(classifier_LR_tune, parameters, cv=2,
                   scoring='f1_micro',return_train_score=True,n_jobs=-1,verbose=2)

clf.fit(x_train_multilabel, y_train)

results = pd.DataFrame.from_dict(clf.cv_results_)

#https://stackoverflow.com/questions/57136676/sklearn-model-selection-gridsearchcv

train_f1= results['mean_train_score']
train_f1_std= results['std_train_score']
cv_f1 = results['mean_test_score']
cv_f1_std=results['std_test_score']
alpha= results['param_estimator__alpha']

log_alpha_SGD_LR=[]
for i in alpha:
    x=math.log(10*i)
    log_alpha_SGD_LR.append(x)

plt.figure(figsize=(10,8))
plt.plot(log_alpha_SGD_LR, train_f1, label='Train f1')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, train_auc - train_auc_std,train_auc + train_auc_std,color='red')

plt.plot(log_alpha_SGD_LR, cv_f1, label='CV f1')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='red')

plt.scatter(log_alpha_SGD_LR, train_f1, label='Train f1 points')
plt.scatter(log_alpha_SGD_LR, cv_f1, label='CV f1 points')

plt.legend()
plt.xlabel("log(10*alpha): hyperparameter",fontsize=18)
plt.ylabel("Micro_f1",fontsize=18)

```

```
plt.title("Hyper parameter Vs micro f1 plot",fontsize=20)
plt.grid()
plt.show()

best_alpha=clf.best_params_
print("="*100)
print("Best_hyperparameter_SGD_LR_Bow:",best_alpha)
Best_hyperparameter_SGD_LR_Bow=best_alpha.get("alpha")
print("="*100)
results.head(2)

print("Time taken to run this cell :", datetime.now() - start)
```

Fitting 2 folds for each of 4 candidates, totalling 8 fits

```
[CV] estimator__alpha=1e-05 .....
[CV] estimator__alpha=1e-05 .....
[CV] estimator__alpha=0.0001 .....
[CV] estimator__alpha=0.0001 .....
[CV] estimator__alpha=0.001 .....
[CV] estimator__alpha=0.001 .....
[CV] estimator__alpha=0.001 .....
[CV] estimator__alpha=0.01 .....
[CV] estimator__alpha=0.01 .....
[CV] estimator__alpha=0.001, score=0.41349942293150793, total= 6.4min
```

```
[Parallel(n_jobs=-1)]: Done 1 tasks | elapsed: 6.6min
```

```
[CV] .. estimator__alpha=0.01, score=0.2607410798830421, total= 6.5min
```

```
[Parallel(n_jobs=-1)]: Done 2 out of 8 | elapsed: 6.7min remaining: 20.1min
```

```
[CV] estimator__alpha=0.0001, score=0.3304578872927577, total= 6.6min
```

```
[Parallel(n_jobs=-1)]: Done 3 out of 8 | elapsed: 6.7min remaining: 11.2min
```

```
[CV] . estimator__alpha=0.001, score=0.3960370129578064, total= 6.6min
```

```
[Parallel(n_jobs=-1)]: Done 4 out of 8 | elapsed: 6.8min remaining: 6.8min
```

```
[CV] . estimator__alpha=0.01, score=0.23696758885996638, total= 6.7min
```

```
[Parallel(n_jobs=-1)]: Done 5 out of 8 | elapsed: 6.8min remaining: 4.1min
```

```
[CV] estimator__alpha=0.0001, score=0.3512897624251587, total= 6.8min
```


```
[Parallel(n_jobs=-1)]: Done 6 out of 8 | elapsed: 6.9min remaining: 2.3min
```

```
[CV] estimator__alpha=1e-05, score=0.27927491077582767, total= 7.3min
```

```
[CV] . estimator__alpha=1e-05, score=0.3182101881818247, total= 7.3min
```

```
[Parallel(n_jobs=-1)]: Done 8 out of 8 | elapsed: 7.4min remaining: 0.0s
```

```
[Parallel(n_jobs=-1)]: Done 8 out of 8 | elapsed: 7.4min finished
```



```

=====
=====
Best_hyperparameter_SGD_LR_BoW: {'estimator__alpha': 0.001}
=====
=====
Time taken to run this cell : 0:12:38.403597

```

2.2: SGDClassifier(Log Loss)/Logistic Regression with Best Alpha

```

In [30]: start = datetime.now()
classifier = OneVsRestClassifier(SGDClassifier(loss='log',
                                              alpha=0.001, penalty='l1'),
                               n_jobs=-1)

classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict (x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, r

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, r

print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.2453
Hamming loss  0.00269525
Micro-average quality numbers
Precision: 0.7557, Recall: 0.5416, F1-measure: 0.6310
Macro-average quality numbers
Precision: 0.2569, Recall: 0.1666, F1-measure: 0.1767

```

	precision	recall	f1-score	support
0	0.98	0.99	0.98	36915
1	0.22	0.06	0.09	140
2	0.03	0.16	0.06	37
3	0.38	0.16	0.22	4486
4	0.47	0.27	0.34	784
5	0.78	0.51	0.62	486
6	0.35	0.62	0.45	220
7	0.11	0.06	0.08	33
8	0.04	0.14	0.06	7
9	0.35	0.18	0.24	44
10	0.50	0.25	0.34	244
11	0.22	0.22	0.22	255

3: Assignment Task 3---->> Try OneVsRestClassifier with Linear-SVM (SGDClassifier with loss-hinge)

3.1: Try OneVsRestClassifier with Linear-SVM (SGDClassifier with loss-hinge)

```
In [31]: start = datetime.now()

classifier = OneVsRestClassifier(SGDClassifier(loss='hinge',
                                             alpha=0.00001, penalty='l1'),
                               n_jobs=-1)

classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)

print("Accuracy :", metrics.accuracy_score(y_test, predictions))
print("Hamming loss ", metrics.hamming_loss(y_test, predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)
```

```
Accuracy : 0.16905
Hamming loss  0.0040626
Micro-average quality numbers
Precision: 0.5191, Recall: 0.6124, F1-measure: 0.5619
Macro-average quality numbers
Precision: 0.1698, Recall: 0.2492, F1-measure: 0.1877
```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	36915
1	0.13	0.14	0.14	140
2	0.17	0.19	0.18	37
3	0.22	0.23	0.22	4486
4	0.36	0.35	0.36	784
5	0.60	0.59	0.59	486
6	0.45	0.45	0.45	220
7	0.05	0.09	0.07	33
8	0.04	0.29	0.08	7
9	0.10	0.20	0.14	44
10	0.36	0.54	0.43	244
11	0.15	0.22	0.18	255

4: CONCLUSION

```
In [11]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Sr.No", "MODEL", "Vectorizer", "HYPERPARAMETER", 'MICRO_F1_SCORE']

x.add_row(["1", "SGD(LR)Hyperparameter_tuned", "BOW((1,4)grams)", "0.001", "0.631"])
x.add_row(["2", "SGD(linearSVM)", "BOW((1,4)grams)", "0.00001", "0.561"])
print(x)
```

```
+-----+-----+-----+-----+-----+
-----+
| Sr.No |          MODEL          | Vectorizer | HYPERPARAMETER | MICRO_F1_SCORE |
+-----+-----+-----+-----+-----+
|  1    | SGD(LR)Hyperparameter_tuned | BOW((1,4)grams) | 0.001          | 0.631          |
|  2    |          SGD(linearSVM)      | BOW((1,4)grams) | 0.00001        | 0.561          |
+-----+-----+-----+-----+-----+
```